

Task DL. The Colruyt Deep Learning Challenge

- Student name: Yori Verbist
- Number: R0697651

Last year, you probably heard about a pilot project of Colruyt, one of the largest supermarket chains in Belgium. Since August 2019, Colruyt has been using Deep Learning in two of its supermarkets to recognize fruit and vegetables and weight them automatically. Now it's up to you to build and train such an image recognition system yourself.

We are also going to do a small competition. Try to achieve an accuracy as high as possible. During the last lesson, we will give you 20 images of fruits. The student who can classify the most images correctly wins the competition. In case of an ex aequo, the student with the highest accuracy wins. Eternal fame will be his part.

You can read and watch the news item on vrt via the link below (in Dutch).

[vrt NWS](#)

Colruyt made a little video which you can watch below.

```
In [1]: %%HTML
<iframe width="640" height="360" src="https://www.youtube.com/embed/RPjt1-
```

Deep learning in Colruyt Kortrijk



1. The dataset

First you need a large amount of train and test images of fruits. I've downloaded about 500 images of 10 different fruits from <https://www.flickr.com/>. You can find them in this rar-file:

Colruyt

Start by unpacking the training/validation dataset. First have a look at the different folders and images.

2. Binary classification

Next, build and train a classifier (like we did with the cats and dogs example). At first make the classification task a little bit simpler with only two fruits (binary = two-class). This will result in lower complexity and hopefully faster training on your laptop.

What two fruits did you select? apples and peaches because they look alike

Insert a screenshot of the training process and the final accuracy achieved.

```
In [29]: from fastai.vision.all import *
        from pathlib import Path
```

```
In [64]: # Insert the code for building and training your classifier.
        path1 = Path("binary")

        fruits1 = DataBlock(
            blocks=(ImageBlock, CategoryBlock),
            get_items=get_image_files,
            splitter=RandomSplitter(valid_pct=0.2, seed=42),
            get_y=parent_label,
            item_tfms=RandomResizedCrop(224, min_scale=0.5),
            batch_tfms=aug_transforms())

        dls = fruits1.dataloaders(path1)
        #dls.valid.show_batch(max_n=16, nrows=4)

        learn1 = cnn_learner(dls, resnet18, metrics=accuracy)
        learn1.fine_tune(5)
```

epoch	train_loss	valid_loss	accuracy	time
-------	------------	------------	----------	------

0	0.893534	0.264972	0.896341	00:04
---	----------	----------	----------	-------

epoch	train_loss	valid_loss	accuracy	time
-------	------------	------------	----------	------

0	0.312991	0.127075	0.951219	00:05
---	----------	----------	----------	-------

1	0.270857	0.180389	0.939024	00:05
---	----------	----------	----------	-------

2	0.234716	0.103726	0.963415	00:05
---	----------	----------	----------	-------


3	0.182895	0.095573	0.969512	00:05
---	----------	----------	----------	-------

4	0.148316	0.097250	0.969512	00:05
---	----------	----------	----------	-------

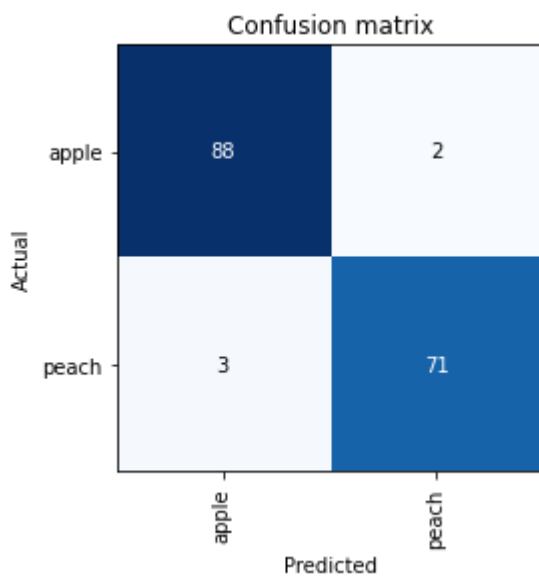
```
In [69]: learn1.recorder.values[-1][-1]
```

```
Out[69]: 0.9695122241973877
```

```
In [70]: dls.train.show_batch(max_n=16, nrows=4)
```



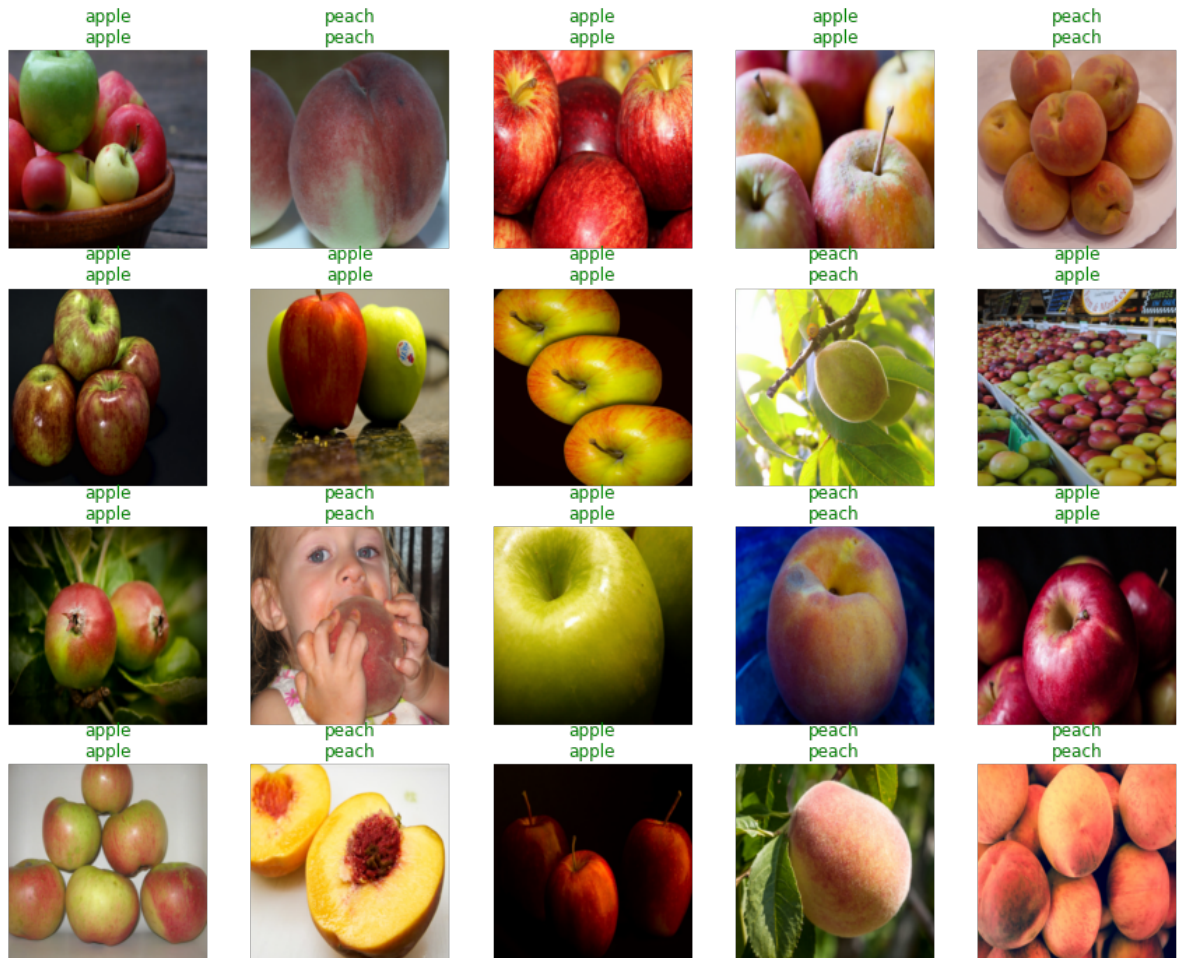
```
In [71]: interp1 = ClassificationInterpretation.from_learner(learn1)
interp1.plot_confusion_matrix()
```



```
In [72]: interp1.plot_top_losses(5, nrows=1)
```



In [73]: `learn1.show_results(max_n=20)`



Download at least ten images of the two fruits (five of each) you selected and see if your model can classify them correctly.

3. Classification with all fruits

Now we make things a little more complicated and challenging. Repeat the proces from above with all the (10) classes. It might take some time now to train the model! You also need to change your program because there are 10 classes now.

Once you're finished, try to achieve an accuracy as high as possible. This is not an exact science. So you will have to experiment with the parameters of your CNN. Possible things to consider:

- the input shape of the images
- the number of filters

- the use of layers to prevent overfitting
- the number of layers
- the number of epochs
- the number of steps per epoch
- ...

It might be a good idea to have a look at some existing famous CNN Architectures for Image Classification (VGGNet, ...), built by researchers around the world. So search the Internet and find some good examples of CNN's and look for some best practices.

This website can be a good start: [Five Powerful CNN Architectures](#)

Write a small report (10 lines) about the things you've tried to optimize your classifier and which configuration gave the best results. Where did you find the information?

Insert a screenshot of the training process and the final accuracy achieved.

```
In [40]: from fastai.vision.all import *
        from pathlib import Path
```

```
In [49]: # Insert the code for building and training your classifier.
        path = Path("all")

        fruits = DataBlock(
            blocks=(ImageBlock, CategoryBlock),
            get_items=get_image_files,
            splitter=RandomSplitter(valid_pct=0.2, seed=42),
            get_y=parent_label,
            item_tfms=RandomResizedCrop(224, min_scale=0.5),
            batch_tfms=aug_transforms())

        dls = fruits.dataloaders(path)
        #dls.valid.show_batch(max_n=16, nrows=4)

        learn = cnn_learner(dls, resnet18, metrics=accuracy)
        learn.fine_tune(20)
```

epoch	train_loss	valid_loss	accuracy	time
-------	------------	------------	----------	------

0	1.143781	0.275712	0.919717	00:11
---	----------	----------	----------	-------

epoch	train_loss	valid_loss	accuracy	time
-------	------------	------------	----------	------

0	0.340018	0.231448	0.935065	00:12
---	----------	----------	----------	-------

1	0.252680	0.202523	0.946871	00:12
---	----------	----------	----------	-------

2	0.212636	0.199563	0.950413	00:12
---	----------	----------	----------	-------

3	0.163889	0.199313	0.949233	00:12
---	----------	----------	----------	-------

4	0.134962	0.189700	0.952775	00:12
---	----------	----------	----------	-------

5	0.111481	0.198960	0.956316	00:12
---	----------	----------	----------	-------

6	0.095298	0.172625	0.957497	00:12
---	----------	----------	----------	-------



7	0.076800	0.201236	0.953955	00:12
---	----------	----------	----------	-------

epoch	train_loss	valid_loss	accuracy	time
8	0.066619	0.187142	0.956316	00:12
9	0.064351	0.238715	0.953955	00:12
10	0.058398	0.219932	0.958678	00:12
11	0.045188	0.169479	0.959858	00:12
12	0.032612	0.163163	0.962220	00:12
13	0.031780	0.161692	0.957497	00:12
14	0.026243	0.176382	0.962220	00:12
15	0.018629	0.172599	0.964581	00:12
16	0.016365	0.174691	0.963400	00:12
17	0.017444	0.166571	0.963400	00:12
18	0.012222	0.177221	0.961020	00:12

In [60]: `learn.recorder.values[-1][1]`

Out[60]: 0.9622195959091187

In [59]: `interp = ClassificationInterpretation.from_learner(learn)`
`interp.plot_confusion_matrix()`
`interp.plot_top_losses(5, nrows=1)`



In [51]: `learn.show_results(max_n=10)`



4. Classify the sample images

Now that your classifier is ready, have a look at the images in the `single_images` folder. Write a Python program that automatically retrieves a list of all the images in this folder and classifies these images. Print a list with in the first column the name of the image file and in the second column the recognized fruit. Check how many fruits were classified correctly. Are you ready to compete in our Colruyt Deep Learning Challenge?

```
In [52]: # Insert a list of images and classification
# Correct x/20 ?
for i in range(20):
    name = "colruyt_orig/single_images/img" + f'{i + 1:02}' + ".jpg"
    x = name.find("img")
    predict = learn.predict(name)
    print(name[x:] + "\t\t", predict[0] + "\t certainty:", predict[2][pred
```

```
img01.jpg      apple  certainty: tensor(1.0000)
img02.jpg      cherry certainty: tensor(1.0000)
img03.jpg      apple  certainty: tensor(1.0000)
img04.jpg      cherry certainty: tensor(1.0000)
img05.jpg      kiwi   certainty: tensor(1.)
img06.jpg      peach  certainty: tensor(1.0000)
img07.jpg      fig    certainty: tensor(1.0000)
img08.jpg      rambutan      certainty: tensor(1.0000)
img09.jpg      grape  certainty: tensor(1.0000)
img10.jpg      lemon  certainty: tensor(1.0000)
```

4. Task

file:///home/iori/Documents/School/3ITF/AI/AI_...

img11.jpg	grape	certainty: tensor(1.)
img12.jpg	strawberry	certainty: tensor(1.0000)
img13.jpg	lemon	certainty: tensor(1.0000)
img14.jpg	strawberry	certainty: tensor(1.)
img15.jpg	kiwi	certainty: tensor(1.)
img16.jpg	pear	certainty: tensor(1.0000)
img17.jpg	pear	certainty: tensor(1.0000)
img18.jpg	peach	certainty: tensor(1.)
img19.jpg	fig	certainty: tensor(1.0000)

Print your Jupyter Notebook to pdf and upload it via Canvas.

In []: